

Mohit Kumar

✉ 5102mohit@gmail.com ☎ +91-8374106416 📁 Portfolio in LinkedIn 🐙 Github

Experience

Associate Data Scientist

Blackcoffer

Delhi, India

August 2025 – Present

- Engineered a full stack AI recruiting platform with FastAPI & Langflow, automating lead management and outreach, reducing recruiter manual effort by 70%.
- Orchestrated a multi-tool agent in Langflow, leveraging the Model Context Protocol (MCP) to integrate Zoho CRM and Salesmsg APIs for automated lead management and SMS outreach.
- Architected and containerized multi-service application with Docker, deployed to Azure VM with Azure Database for PostgreSQL and Redis for scalable data handling and caching.

Machine Learning Intern

Labmentix

Bangalore, India

May 2025 – July 2025

- Developed a machine learning model to automatically classify support tickets by issue type and urgency, reducing manual triage time.
- Implemented a rule-based entity extractor for date recognition and keyword matching for product and complaint identification.
- Built an automated data pipeline with Python and Selenium to extract and structure key information from more than 20,000+ financial transcripts.

Projects

Legal Document Summarizer

- Fine-tuned a TinyLlama-1.1B model using QLoRA to create a specialist legal summarization agent, enhancing domain-specific understanding and summary precision.
- Orchestrated the fine-tuned agent with the Gemini API in a hybrid RAG architecture, enabling stateful, conversational Q&A and efficient retrieval from summarized legal documents.

ChromaTune - Multimodal Music Recommender

- Enhanced a multimodal RAG system that generates Spotify song recommendations by interpreting the vibe from user-provided images and videos.
- Optimized a data pipeline to extract audio embeddings, generate semantic song descriptions from multimodal inputs, and index them into an in-memory FAISS vector store for fast, high-precision similarity search.

Technical Skills

Languages: Python, SQL, C/C++, JavaScript, HTML/CSS, R

Machine Learning & Generative AI: Pandas, NumPy, Scikit-Learn, PyTorch, Hugging Face, Langflow, Fine-tuning (LoRA, QLoRA), Retrieval-Augmented Generation (RAG), Vector Databases (FAISS, ChromaDB)

Backend & Databases: FastAPI, Django, MongoDB, PostgreSQL, MySQL, Neo4j

Cloud & DevOps: Azure, Docker, MLOps, Grafana

Education

SRM University

Bachelor of Technology - Computer Science and Engineering

Sept 2021 – June 2025

Achievements & Extracurricular

Zonal chess player: Represented university as a zonal-level chess player, showcasing strategic and analytical skills. Achieved multiple victories in university chess tournaments, demonstrating consistency and excellence.

DeepLearning.AI: Completed DeepLearning.AI specialization on Generative AI with Large Language Models.